

# 1 Theoretical Analysis of GME Exploration Bonuses

## 1.1 Background: Linear MDP and LSVI-UCB

In linear MDPs, transition kernels and the reward function are assumed to be linear. ? formalizes the definition of linear MDPs as follows:

**Definition 1.1** (Linear MDP). *An MDP  $(\mathcal{S}, \mathcal{A}, H, \mathcal{P}, r)$  is a linear MDP with a feature map  $\phi : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}^d$  if for any  $h \in [H]$ , there exist  $d$  unknown (signed) measures  $\mu_h = (\mu_h^{(1)}, \dots, \mu_h^{(d)})$  over  $\mathcal{S}$  and an unknown vector  $\theta_h \in \mathbb{R}^d$ , such that for any  $(x, a) \in \mathcal{S} \times \mathcal{A}$ :*

$$\mathcal{P}_h(\cdot|x, a) = \langle \phi(x, a), \mu_h(\cdot) \rangle, \quad r_h(x, a) = \langle \phi(x, a), \theta_h \rangle. \quad (1)$$

In linear MDPs, LSVI-UCB achieves near-optimal worst-case regret. The key idea of LSVI-UCB is to use optimistic  $Q$ -values obtained by adding an UCB bonus  $r^{\text{ucb}}$  to the estimated  $Q$ -values. The UCB bonus is defined as:

$$r_t^{\text{ucb}} = \beta \cdot [\eta(s_t, a_t)^\top \Lambda_t^{-1} \eta(s_t, a_t)]^{1/2},$$

where  $\beta$  is a constant,  $\Lambda_t = \sum_{i=0}^m \eta(x_t^i, a_t^i) \eta(x_t^i, a_t^i)^\top + \lambda \cdot \mathbf{I}$  is the Gram matrix, and  $m$  is the index of the current episode. The UCB bonus measures the epistemic uncertainty of state-action pairs and has been proven to be efficient. The LSVI-UCB algorithm is described in Algorithm 1. Each iteration of LSVI-UCB consists of two parts: first, in lines 3-6, the agent executes a policy based on  $Q_t$ ; second, in lines 7-11, the  $Q$ -function parameters  $\chi_t$  are updated via regularized least squares:

$$\chi_t \leftarrow \arg \min_{\chi \in \mathbb{R}^d} \sum_{i=0}^m \left[ r_t(s_t^i, a_t^i) + \max_{a \in \mathcal{A}} Q_{t+1}(s_{t+1}^i, a) - \chi^\top \eta(s_t^i, a_t^i) \right]^2 + \lambda \|\chi\|^2, \quad (2)$$

where  $m$  is the number of episodes and  $i$  is the episode index. This least squares problem has a closed-form solution:

$$\chi_t = \Lambda_t^{-1} \sum_{\tau=0}^m \eta(x_t^\tau, a_t^\tau) \left[ r_t(x_t^\tau, a_t^\tau) + \max_a Q_{t+1}(x_{t+1}^\tau, a) \right],$$

where  $\Lambda_t$  is the Gram matrix. The action-value function is estimated via  $Q_t(s, a) \approx \chi_t^\top \eta(s, a)$ .

LSVI-UCB constructs confidence intervals for the  $Q$ -function using the UCB bonus (line 10):  $r^{\text{ucb}} = \beta [\eta(s, a)^\top \Lambda_t^{-1} \eta(s, a)]^{1/2}$ , which measures the epistemic uncertainty of state-action pairs. Theoretical analysis shows that with appropriate choices of  $\beta$  and  $\lambda$ , LSVI-UCB achieves near-optimal worst-case regret  $\tilde{\mathcal{O}}(\sqrt{d^3 T^3 L^3})$ , where  $L$  is the total number of steps. Next, we establish a theoretical connection between the exploration bonus in GME and the UCB bonus.

---

**Algorithm 1** LSVI-UCB Algorithm for Linear MDPs

---

```

1: Initialize:  $\Lambda_t \leftarrow \lambda \cdot \mathbf{I}$  and  $w_h \leftarrow 0$ 
2: for episode  $m = 0$  to  $M - 1$  do
3:   Receive initial state  $s_0$ 
4:   for step  $t = 0$  to  $T - 1$  do
5:     Execute action  $a_t = \arg \max_{a \in \mathcal{A}} Q_t(s_t, a)$  and observe  $s_{t+1}$ 
6:   end for
7:   for step  $t = T - 1$  downto  $0$  do
8:      $\Lambda_t \leftarrow \sum_{i=0}^m \eta(x_t^i, a_t^i) \eta(x_t^i, a_t^i)^\top + \lambda \cdot \mathbf{I}$ 
9:      $\chi_t \leftarrow \Lambda_t^{-1} \sum_{i=0}^m \eta(x_t^i, a_t^i) [r_t(x_t^i, a_t^i) + \max_a Q_{t+1}(x_{t+1}^i, a)]$ 
10:     $Q_t(\cdot, \cdot) = \min \left\{ \chi_t^\top \eta(\cdot, \cdot) + \alpha [\eta(\cdot, \cdot)^\top \Lambda_t^{-1} \eta(\cdot, \cdot)]^{1/2}, T \right\}$ 
11:   end for
12: end for

```

---

## 1.2 Theoretical Connection Between GME and LSVI-UCB

In linear MDPs, we represent the prior model as a linear combination of state-action encodings, i.e.,  $s_{t+1} = W^\top \phi(s_t, a_t) + \epsilon_t$ , where  $\epsilon_t \sim \mathcal{N}(0, \sigma^2 I)$ , and we assume the parameters follow a prior distribution  $W \sim \mathcal{N}(0, \Lambda_0^{-1})$ . We use standard Bayesian analysis to illustrate our conclusions.

The prior distribution of the parameter matrix  $W \in \mathbb{R}^{d \times d}$  is:

$$p(W) = \mathcal{N}(W | \mathbf{0}, \Lambda_0^{-1})$$

where  $\Lambda_0 = \lambda I$  is the prior variance matrix. Based on the conjugate prior property of Gaussian distributions, the posterior distribution after  $t$  observations is updated as follows:

$$\begin{aligned}\Lambda_t &= \sum_{i=1}^t \phi_i \phi_i^\top + \Lambda_0 \quad (\text{Variance Matrix Update}) \\ \hat{W}_t &= \Lambda_t^{-1} \left( \sum_{i=1}^t \phi_i s_{i+1}^\top \right) \quad (\text{Mean Matrix Update})\end{aligned}$$

where  $\phi_i = \phi(s_i, a_i) \in \mathbb{R}^d$  is the state-action feature vector,  $s_{i+1} \in \mathbb{R}^d$  is the next-state observation,  $\Lambda_t \in \mathbb{R}^{d \times d}$  is the posterior variance matrix, and  $\hat{W}_t \in \mathbb{R}^{d \times d}$  is the posterior mean matrix.

**Theorem 1.2** (Equivalence Between GME Exploration Bonus and UCB). *In linear MDPs, assuming the parameter matrix  $W$  follows a Gaussian prior  $W \sim \mathcal{N}(0, \Lambda_0^{-1})$ , and the latent state  $z_t = W\phi(s_t, a_t)$  follows a Gaussian distribution  $z_t | \mathcal{D}_t \sim \mathcal{N}(0, \phi_t^\top \Lambda_t^{-1} \phi_t I)$ , there exist constants  $\beta_1, \beta_2 > 0$  such that the GME exploration bonus satisfies:*

$$\beta_1 \cdot \sqrt{\phi_t^\top \Lambda_t^{-1} \phi_t} \leq r_t^{\text{GME}} \leq \beta_2 \cdot \sqrt{\phi_t^\top \Lambda_t^{-1} \phi_t} \quad (3)$$

where  $\phi_t = \phi(s_t, a_t)$  and  $\Lambda_t = \lambda I + \sum_{i=1}^t \phi_i \phi_i^\top$ .

*Proof.* The exploration bonus in GME is formulated as:

$$r_t^{\text{GME}} = \mathcal{H}[p(z_t | \mathcal{D}_t)] + D_{\text{KL}}[p(z_t | \mathcal{D}_t) \| q(z_t | s_t)]$$

Let  $W$  follow the matrix normal distribution  $\mathcal{MN}(0, \Lambda_t^{-1}, I)$  with vectorization  $\text{vec}(W) \sim \mathcal{N}(0, \Lambda_t^{-1} \otimes I)$ . The latent variable  $z_t = W\phi_t$  has covariance:

$$\begin{aligned}\Sigma_t &= \mathbb{E}[(z_t - \mathbb{E}z_t)(z_t - \mathbb{E}z_t)^\top] \\ &= \phi_t^\top \mathbb{E}[WW^\top] \phi_t \cdot I \\ &= \phi_t^\top \Lambda_t^{-1} \phi_t \cdot I \quad (\text{by } \mathbb{E}[WW^\top] = \Lambda_t^{-1})\end{aligned} \quad (4)$$

For  $p(z_t | \mathcal{D}_t) = \mathcal{N}(0, \phi_t^\top \Lambda_t^{-1} \phi_t \cdot I)$ :

$$\begin{aligned}\mathcal{H}[p(z_t | \mathcal{D}_t)] &= \frac{d}{2} \log(2\pi e) + \frac{1}{2} \log \det(\Sigma_t) \\ &= \frac{d}{2} \log(2\pi e) + \frac{d}{2} \log(\phi_t^\top \Lambda_t^{-1} \phi_t)\end{aligned} \quad (5)$$

The covariance matrix is calculated as follows:

$$\begin{aligned}
\Sigma_t &= \mathbb{V}[W\phi_t] \\
&= \mathbb{E}[(W\phi_t - \mathbb{E}W\phi_t)(W\phi_t - \mathbb{E}W\phi_t)^\top] \\
&= \phi_t^\top \mathbb{E}[(W - \hat{W}_t)(W - \hat{W}_t)^\top] \phi_t \\
&= \phi_t^\top \left( \mathbb{E}[WW^\top] - \hat{W}_t \hat{W}_t^\top \right) \phi_t \\
&= \phi_t^\top \Lambda_t^{-1} \phi_t \cdot I \quad (\text{Based on posterior covariance } \mathbb{V}[W] = \Lambda_t^{-1}) \tag{6}
\end{aligned}$$

Let  $q(z_t|s_t) = \mathcal{N}(0, \phi_t^\top \Lambda_0^{-1} \phi_t \cdot I)$ . The KL divergence is:

$$\begin{aligned}
D_{\text{KL}}[p||q] &= \frac{1}{2} \left[ \log \frac{\det \Sigma_q}{\det \Sigma_p} + \text{tr}(\Sigma_q^{-1} \Sigma_p) - d \right] \\
&= \frac{1}{2} \left[ d \log \frac{\phi_t^\top \Lambda_0^{-1} \phi_t}{\phi_t^\top \Lambda_t^{-1} \phi_t} + \frac{\phi_t^\top \Lambda_t^{-1} \phi_t}{\phi_t^\top \Lambda_0^{-1} \phi_t} \cdot d - d \right] \tag{8}
\end{aligned}$$

Combining the two terms, we have:

$$r_t^{\text{GME}} = \frac{d}{2} \log(\phi_t^\top \Lambda_t^{-1} \phi_t) + \frac{d}{2} \left[ \frac{\phi_t^\top \Lambda_t^{-1} \phi_t}{\phi_t^\top \Lambda_0^{-1} \phi_t} - \log \frac{\phi_t^\top \Lambda_t^{-1} \phi_t}{\phi_t^\top \Lambda_0^{-1} \phi_t} - 1 \right] + C \tag{9}$$

where  $C = \frac{d}{2} [\log(2\pi e) + \log(\phi_t^\top \Lambda_0^{-1} \phi_t)]$ .

Now we let  $v_t = \sqrt{\phi_t^\top \Lambda_t^{-1} \phi_t}$  and note that  $\lambda_{\min}(\Lambda_t) \geq \lambda$ , then:

$$\frac{1}{\lambda} \|\phi_t\|^2 \geq \phi_t^\top \Lambda_t^{-1} \phi_t \geq \frac{\|\phi_t\|^2}{\lambda + t} \geq 0 \quad (\text{by } \Lambda_t \preceq (\lambda + t)I) \tag{10}$$

Define  $f(x) = \frac{d}{2} \log x + \frac{d}{2} \left[ \frac{x}{c} - \log \frac{x}{c} - 1 \right]$  where  $c = \phi_t^\top \Lambda_0^{-1} \phi_t$ . Using inequalities:

$$\frac{x}{2} \leq \log(1+x) \leq x \quad \text{for } 0 \leq x \leq 1 \tag{11}$$

$$\log x \leq x - 1 \quad \forall x > 0 \tag{12}$$

We can show  $\exists \beta_1, \beta_2 > 0$  such that:

$$\beta_1 v_t \leq \sqrt{f(v_t^2) + C} \leq \beta_2 v_t \quad (\text{via case analysis on } v_t \leq 1 \text{ and } v_t > 1) \tag{13}$$

Specifically, choosing:

$$\beta_1 = \sqrt{\frac{d}{2} \min \left\{ 1, \frac{1}{\lambda} \right\}} \quad (14)$$

$$\beta_2 = \sqrt{d \left( 1 + \frac{1}{\lambda} \right)} \quad (15)$$

satisfies the inequality for all  $t \geq 0$ . □

## 2 Atari100k and Atari1000k

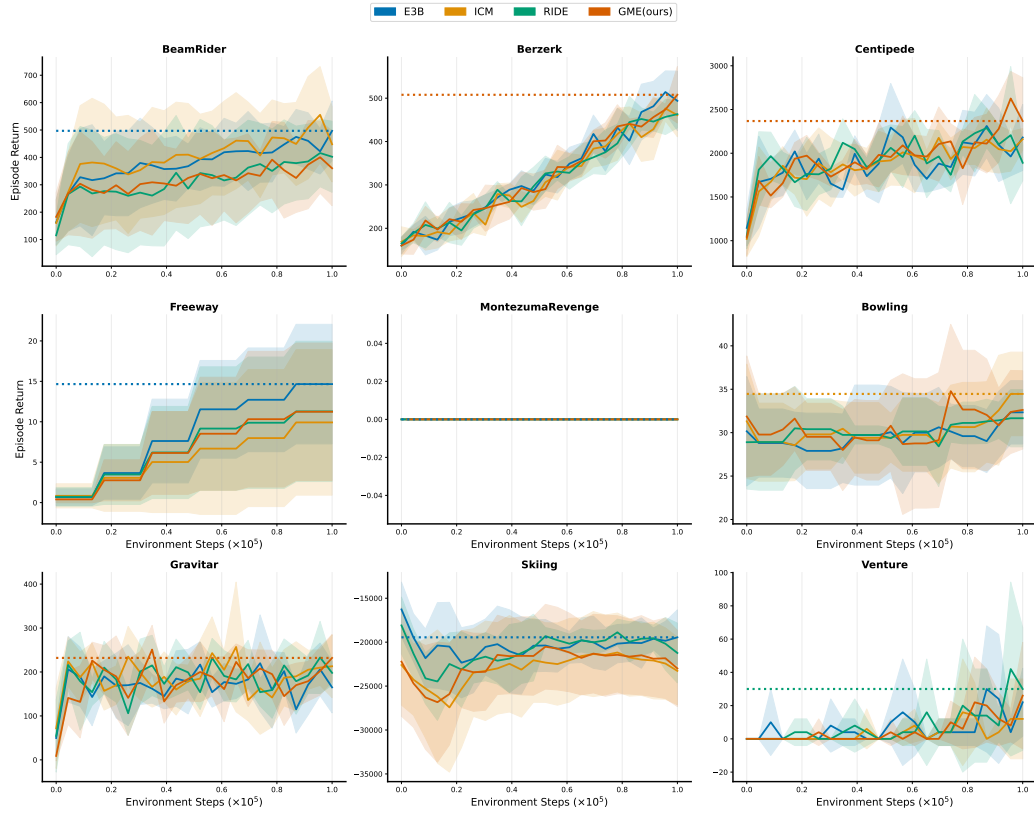


Figure 1: Results on 9 Atari games with 100k step

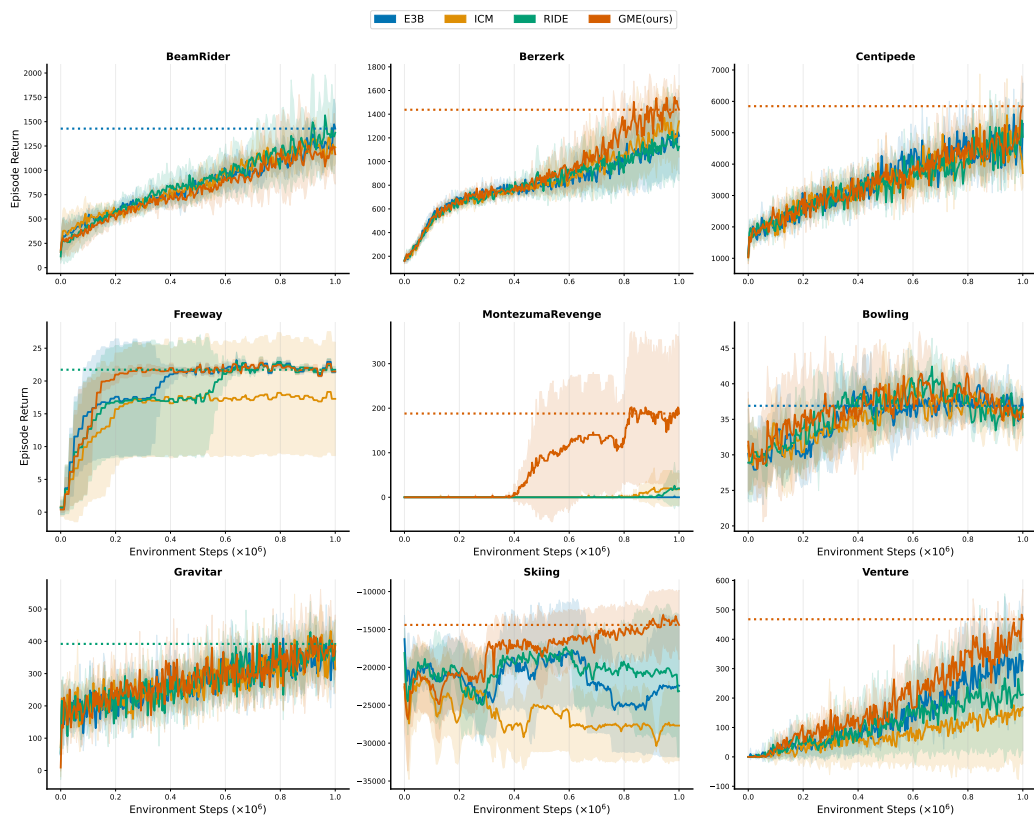


Figure 2: Results on 9 Atari games with 1000k step